



Solution Set

Question 1

Consider a Brownian risk factor used to drive an equity or rates model. Define the state variables (filtration, Brownian motion, adapted processes, quadratic variation, and stopping times) and derive the Ito dynamics, martingale condition, and exponential density process. State the measure, filtration, and tradability assumptions.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

Answer key

Define the tradable state variables first: filtration, Brownian motion, adapted processes, quadratic variation, and stopping times. The payoff is then written at the relevant horizon and valued under the pricing measure.

The derivation uses the economic relation for Ito dynamics, martingale condition, and exponential density process. If a hedge or replication argument is used, the stochastic exposure is neutralized locally and the remaining cash account earns the risk-free or collateral rate.

The final expression must agree with the payoff at maturity and with the obvious limiting cases.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt \\ d\Pi = dV - \Delta dS, \quad d\Pi = r\Pi dt$$

Question 2

For a Brownian risk factor used to drive an equity or rates model, explain which probability measure is natural for pricing. Write the density-process logic needed if the starting model is under the physical measure.

Answer key

The pricing measure is chosen so discounted tradable prices are martingales under the chosen numeraire.

The Radon-Nikodym density reweights probabilities; it does not change the contractual payoff. In Brownian models, Girsanov changes the drift and leaves the quadratic variation unchanged.

A useful check is that the density has expectation one and that the resulting discounted price process has zero drift.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt \\ Z_t = \frac{dQ}{dP} \Big|_{\mathcal{F}_t}, \quad E^Q[X | \mathcal{F}_t] = \frac{1}{Z_t} E^P[Z_T X | \mathcal{F}_t]$$

Question 3

Identify the discounted tradable or density-adjusted process in a Brownian risk factor used to drive an equity or rates model. Prove the martingale condition and name one modelling choice that would break it.

Answer key

The martingale candidate is the discounted tradable price, or the density-adjusted process after a measure change.

Condition on the earlier information set and verify that the conditional expectation equals the earlier value. Any remaining predictable drift indicates a missing funding, dividend, numeraire, or model adjustment.

For this topic, the practical issue is using future path information when defining a trading rule.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt \\ E[M_t | \mathcal{F}_s] = M_s, \quad 0 \leq s < t$$

Question 4

Using a one-year Brownian path with monthly time steps and volatility 22 percent, construct a two-step numerical calculation. Show intermediate state values before discounting or aggregating.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

Answer key

Use the stated market inputs: a one-year Brownian path with monthly time steps and volatility 22 percent. Compute future state values first, then take the correct conditional average, then discount.

Keep a calculation table with states, probabilities, payoffs, discount factors, and final present values. This makes unit mistakes visible.

The answer should include a sign and magnitude check, not only a final number.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$
$$V_t = e^{-r\Delta t}(qV_u + (1-q)V_d)$$

Question 5

For a Brownian risk factor used to drive an equity or rates model, compute and interpret drift removal and martingale exposure under a changed measure. Explain the hedge sign and the unit of each sensitivity.

Answer key

The relevant exposure is drift removal and martingale exposure under a changed measure. Compute the first-order term, then add convexity or cross terms only when the state variables require them.

A hedge has the opposite sign of the exposure. Delta is per unit of underlying, DV01 is per basis point, and vega is per volatility point or percent depending on convention.

After hedging, report the residual risk that remains.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$
$$dV \approx \Delta dS + \frac{1}{2}\Gamma(dS)^2 + \nu d\sigma$$

Question 6

Suppose the desk observes drift and diffusion coefficients inferred from increments. Write the calibration objective, the weights you would use, and the parameter-identifiability checks.

Answer key

The calibration instruments are drift and diffusion coefficients inferred from increments. Write a weighted least-squares objective using market quotes and model quotes.

Weights should reflect liquidity, bid-ask width, and quote reliability. Parameter constraints should prevent impossible prices or unstable extrapolation.

After calibration, inspect residuals and check whether each parameter has a distinct effect on the quotes.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$
$$\hat{\theta} = \arg \min_{\theta} \sum_i w_i (M_i(\theta) - Q_i^{\text{mkt}})^2$$

Question 7

State the terminal and boundary conditions for a Brownian risk factor used to drive an equity or rates model. Explain how stopping times and path barriers defined by the available information prevents economically impossible prices.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

Answer key

The terminal condition comes from the payoff. Boundary conditions come from economic limits such as zero spot, very large spot, barrier touch, exercise, or discount-factor positivity.

For this product, the key boundary idea is stopping times and path barriers defined by the available information.

A boundary condition is wrong if it permits arbitrage, negative values for non-negative payoffs, or a price that violates obvious monotonicity.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$V(T, S) = g(S), \quad V(t, S_{\text{boundary}}) = \text{economic boundary value}$$

Question 8

Design a Monte Carlo or lattice estimator for a Brownian risk factor used to drive an equity or rates model. Specify Euler-Maruyama paths and quadratic-variation diagnostics, the estimator, the discounting rule, and the standard error.

Answer key

Simulate Euler-Maruyama paths and quadratic-variation diagnostics. For each path, compute the payoff using only information available along that path.

The estimator is the discounted sample average. Report standard error and rerun with larger path counts or finer time steps to separate sampling error from discretization error.

For path-dependent products, store the running state variable explicitly.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$\widehat{V}_0 = e^{-rT} \frac{1}{N} \sum_{i=1}^N X^{(i)}, \quad \widehat{\text{SE}} = e^{-rT} \frac{s_X}{\sqrt{N}}$$

Question 9

A junior quant implements a Brownian risk factor used to drive an equity or rates model and gets a stable but wrong answer. Explain how using future path information when defining a trading rule creates the error and design a diagnostic test.

Answer key

The main implementation risk is using future path information when defining a trading rule. This can produce stable but wrong results because the error is systematic.

A good diagnostic compares the implementation with a limiting case or benchmark where the direction of the answer is known.

Change one input at a time and verify monotonicity, units, and convergence.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$\text{error}(N, \Delta t) = \widehat{V}_{N, \Delta t} - V_{\text{benchmark}}$$

Question 10

Write an NumPy algorithm for a Brownian risk factor used to drive an equity or rates model. Include inputs, outputs, diagnostics, and one benchmark case.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

Answer key

Create clear arrays for inputs, states, probabilities, discount factors, and payoffs. Check dimensions before combining them.

Return both the main value and diagnostics such as convergence, residuals, standard error, or benchmark difference.

The routine is acceptable only if the result can be reproduced from the displayed assumptions and data.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$\text{tree cost} = O(n^2), \quad \text{simulation cost} = O(NM)$$

Question 11

Starting from the payoff definition of a Brownian risk factor used to drive an equity or rates model, derive a replication, dynamic-programming, or conditional-expectation representation without skipping the economic argument.

Answer key

Define the tradable state variables first: filtration, Brownian motion, adapted processes, quadratic variation, and stopping times. The payoff is then written at the relevant horizon and valued under the pricing measure.

The derivation uses the economic relation for Ito dynamics, martingale condition, and exponential density process. If a hedge or replication argument is used, the stochastic exposure is neutralized locally and the remaining cash account earns the risk-free or collateral rate.

The final expression must agree with the payoff at maturity and with the obvious limiting cases.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$d\Pi = dV - \Delta dS, \quad d\Pi = r\Pi dt$$

Question 12

Create a stressed-market version of a one-year Brownian path with monthly time steps and volatility 22 percent. Recompute the value or risk metric and explain which term dominates the change.

Answer key

Use the stated market inputs: a one-year Brownian path with monthly time steps and volatility 22 percent. Compute future state values first, then take the correct conditional average, then discount.

Keep a calculation table with states, probabilities, payoffs, discount factors, and final present values. This makes unit mistakes visible.

The answer should include a sign and magnitude check, not only a final number.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$V_t = e^{-r\Delta t}(qV_u + (1-q)V_d)$$

Question 13

Use a second-order Taylor approximation for a Brownian risk factor used to drive an equity or rates model. Separate linear exposure, convexity exposure, cross terms if relevant, and residual model risk.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

Answer key

The relevant exposure is drift removal and martingale exposure under a changed measure. Compute the first-order term, then add convexity or cross terms only when the state variables require them.

A hedge has the opposite sign of the exposure. Delta is per unit of underlying, DV01 is per basis point, and vega is per volatility point or percent depending on convention.

After hedging, report the residual risk that remains.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$dV \approx \Delta dS + \frac{1}{2}\Gamma(dS)^2 + \nu d\sigma$$

Question 14

Show how a change of numeraire or change of measure would simplify one expectation in a Brownian risk factor used to drive an equity or rates model. Identify what becomes a martingale after the change.

Answer key

The pricing measure is chosen so discounted tradable prices are martingales under the chosen numeraire.

The Radon-Nikodym density reweights probabilities; it does not change the contractual payoff. In Brownian models, Girsanov changes the drift and leaves the quadratic variation unchanged.

A useful check is that the density has expectation one and that the resulting discounted price process has zero drift.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$Z_t = \frac{dQ}{dP} \Big|_{\mathcal{F}_t}, \quad E^Q[X | \mathcal{F}_t] = \frac{1}{Z_t} E^P[Z_T X | \mathcal{F}_t]$$

Question 15

Give a proof-style answer showing when the wealth process for a Brownian risk factor used to drive an equity or rates model is self-financing. Explain why predictable trading strategies matter.

Answer key

The martingale candidate is the discounted tradable price, or the density-adjusted process after a measure change.

Condition on the earlier information set and verify that the conditional expectation equals the earlier value. Any remaining predictable drift indicates a missing funding, dividend, numeraire, or model adjustment.

For this topic, the practical issue is using future path information when defining a trading rule.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

$$E[M_t | \mathcal{F}_s] = M_s, \quad 0 \leq s < t$$

Question 16

Market quotes for drift and diffusion coefficients inferred from increments are noisy. Propose a robust calibration workflow with initial guesses, constraints, residual plots, and out-of-sample checks.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

Answer key

The calibration instruments are drift and diffusion coefficients inferred from increments. Write a weighted least-squares objective using market quotes and model quotes.

Weights should reflect liquidity, bid-ask width, and quote reliability. Parameter constraints should prevent impossible prices or unstable extrapolation.

After calibration, inspect residuals and check whether each parameter has a distinct effect on the quotes.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$
$$\hat{\theta} = \arg \min_{\theta} \sum_i w_i (M_i(\theta) - Q_i^{\text{mkt}})^2$$

Question 17

For a Brownian risk factor used to drive an equity or rates model, compare the mathematical boundary condition with the trader's intuition. Give one limiting case where the answer should be obvious.

Answer key

The terminal condition comes from the payoff. Boundary conditions come from economic limits such as zero spot, very large spot, barrier touch, exercise, or discount-factor positivity.

For this product, the key boundary idea is stopping times and path barriers defined by the available information.

A boundary condition is wrong if it permits arbitrage, negative values for non-negative payoffs, or a price that violates obvious monotonicity.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$
$$V(T, S) = g(S), \quad V(t, S_{\text{boundary}}) = \text{economic boundary value}$$

Question 18

Design a convergence study for Euler-Maruyama paths and quadratic-variation diagnostics. Specify path counts, time-step refinement, confidence intervals, and the exact quantity plotted.

Answer key

Simulate Euler-Maruyama paths and quadratic-variation diagnostics. For each path, compute the payoff using only information available along that path.

The estimator is the discounted sample average. Report standard error and rerun with larger path counts or finer time steps to separate sampling error from discretization error.

For path-dependent products, store the running state variable explicitly.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$
$$\hat{V}_0 = e^{-rT} \frac{1}{N} \sum_{i=1}^N X^{(i)}, \quad \widehat{\text{SE}} = e^{-rT} \frac{s_X}{\sqrt{N}}$$

Question 19

Find one way that using future path information when defining a trading rule could pass a superficial unit test. Write a stronger test that would fail for the biased implementation.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$

Answer key

The main implementation risk is using future path information when defining a trading rule. This can produce stable but wrong results because the error is systematic.

A good diagnostic compares the implementation with a limiting case or benchmark where the direction of the answer is known.

Change one input at a time and verify monotonicity, units, and convergence.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$
$$\text{error}(N, \Delta t) = \widehat{V}_{N, \Delta t} - V_{\text{benchmark}}$$

Question 20

Explain a Brownian risk factor used to drive an equity or rates model and code the core routine. Give the derivation outline, pseudocode, and three numerical sanity checks.

Answer key

Create clear arrays for inputs, states, probabilities, discount factors, and payoffs. Check dimensions before combining them.

Return both the main value and diagnostics such as convergence, residuals, standard error, or benchmark difference.

The routine is acceptable only if the result can be reproduced from the displayed assumptions and data.

$$dX_t = \mu(t, X_t)dt + \sigma(t, X_t)dW_t, \quad d\langle W \rangle_t = dt$$
$$\text{tree cost} = O(n^2), \quad \text{simulation cost} = O(NM)$$